

Spark SQL

Шугаев Ильянур
VK.com
ilnur.shug@gmail.com

февраль 2020 г.

История

2006 MapReduce

2009 Hive

2009 Pig

2010 Spark

2013 Shark¹

2015 Impala²

2015 Spark SQL³

¹Reynold S Xin и др. “Shark: SQL and rich analytics at scale”. В: *Proceedings of the 2013 ACM SIGMOD International Conference on Management of data*. 2013, с. 13—24.

²Marcel Kornacker и др. “Impala: A Modern, Open-Source SQL Engine for Hadoop.”. В: *Cidr*. Т. 1. 2015, с. 9.

³Michael Armbrust и др. “Spark sql: Relational data processing in spark”. В: *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. 2015, с. 1383—1394.

Spark

Limitations

- Low-level procedural code
- No optimizations

Shark

- First effort to build a relational interface on Spark
- Shark modified the Apache Hive system to run on Spark

Shark

Limitations

- Shark could only be used to query external data stored in the Hive catalog
- The only way to call Shark from Spark programs was to put together a SQL string
- Hive optimizer was tailored for MapReduce and difficult to extend

Наблюдение

Most data pipelines are combination of relational and procedural algorithms.

Spark SQL — new module in Apache Spark

- DataFrame API
- Catalyst (optimizer)

Определение

DataFrames are collections of structured records that can be manipulated using Spark's procedural API, or using new relational APIs that allow richer optimizations

They can be created directly from Spark's RDDs, enabling relational processing in existing Spark programs.

Spark RDDs

Определение

RDD:

- *Resilient — отказоустойчивый*
- *Distributed — разбитый на партии*
- *Dataset*

read-only, partitioned collection of records

RDDs can be manipulated through operations like map, filter, and reduce, which take functions and ship them to nodes on the cluster.

Spark

Fault-tolerance

- Запомним граф вычислений (linage)
- Тогда если часть данных будет потеряна, то их легко можно восстановить

Spark

Lazy evaluation

- Each RDD represents a “logical plan” to compute a dataset
- Spark waits until action to launch a computation
- Allows to do some simple query optimization, such as pipelining operations (narrow dependencies)

Spark

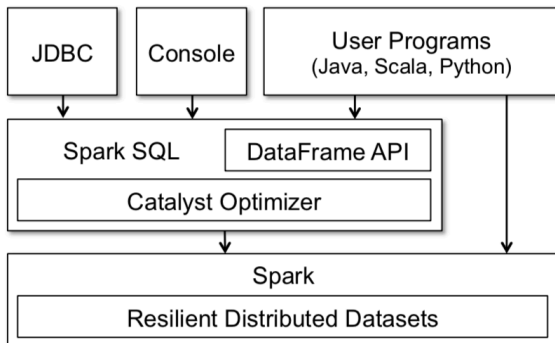
Optimizations

Spark engine does not understand the structure of the data in RDDs (which is arbitrary Java/Python objects) or the semantics of user functions (which contain arbitrary code)

Цели

- 1 Support relational processing both within Spark programs (on native RDDs) and on external data sources using a programmer-friendly API.
- 2 Easily support new data sources, including semi-structured data and external databases amenable to query federation.
- 3 Enable extension with advanced analytics algorithms such as graph processing and machine learning.

PROGRAMMING INTERFACE



- DataFrame — distributed collection of rows with the same *schema*
- DataFrame is equivalent to a table in a relational database
- DataFrames can be manipulated in similar ways to the RDDs
- Schema leads to a more optimized execution

DataFrame

Construction

- From external data sources (HDFS, Hive)
- From existing RDDs (schema inference algorithm)

Замечание

DataFrame can be viewed as an RDD of Row objects, allowing user to call procedural Spark APIs such as map

DataFrame

Execution

- Spark DataFrame are *lazy*
- Spark build *logical plan* before execution
- Laziness enables rich optimization
- Logical plan → Physical plan

DSL

Users can perform relational operations on DataFrames using a domain-specific language (DSL) similar to Python Pandas

Example

Фильмы с наибольшим средним рейтингом

```
1 ratings_df \  
2     .groupby('movie_id') \  
3     .agg(F.mean('rating').alias('mean_rating'),  
4          F.count('rating').alias('ratings_count')) \  
5     .join(movies_df, ratings_df['movie_id'] == movies_df['movieId'],  
6           how='inner') \  
7     .sort(F.col('mean_rating').desc()) \  

```

Difference with native Spark API

- All of these operators build up an abstract syntax tree (AST) of the expression, which is then passed to Catalyst for optimization.
- This is unlike the native Spark API that takes functions containing arbitrary Scala/Java/Python code, which are then opaque to the runtime engine.

DataFrame construction

Schema inference

- While building DataFrame from RDD user can manually define schema
- Spark SQL can automatically infer the schema of the dataset using reflection
- In Python, Spark SQL samples the dataset to perform schema inference due to the dynamic type system

.cache()

- Method cache of DataFrame does the same thing as method persist of RDD
- Caching is particularly useful for interactive queries and for the iterative algorithms common in machine learning

UDF Example

```
1 def get_release_year(title):
2     result = re.match(r'.*(\\d+\\)', title)
3     return int(result.group(1)[1:-1]) if result is not None else None
4
5 get_release_year_udf = F.udf(get_release_year, IntegerType())
6
7 movies_df \
8     .withColumn('year', get_release_year_udf('title')) \
```


Резюме

- The DataFrame API lets developers seamlessly mix procedural and relational methods.

CATALYST

Catalyst contains a general library for representing trees and applying rules to manipulate them

Trees

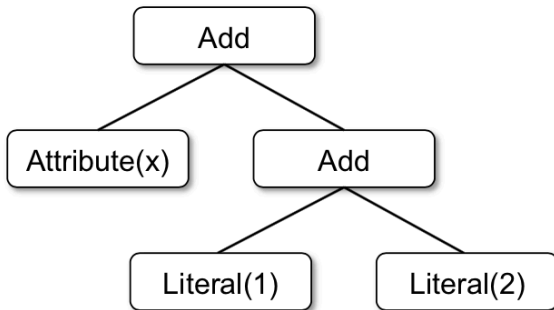


Рис.: Catalyst tree for expression $x + (1 + 2)$

Rules

Определение

Rule: $T \mapsto T'$ — rule maps tree to another tree.

The most common approach is to use a set of pattern matching functions that find and replace subtrees with a specific structure.

Rules

Example

Constant folding

```
1 tree.transform {  
2   case Add(Literal(c1), Literal(c2)) => Literal(c1+c2)  
3   case Add(left, Literal(0)) => left  
4   case Add(Literal(0), right) => right  
5 }
```

Rules

Fixed point

Catalyst groups rules into batches, and executes each batch until it reaches a *fixed point*, that is, until the tree stops changing after applying its rules.

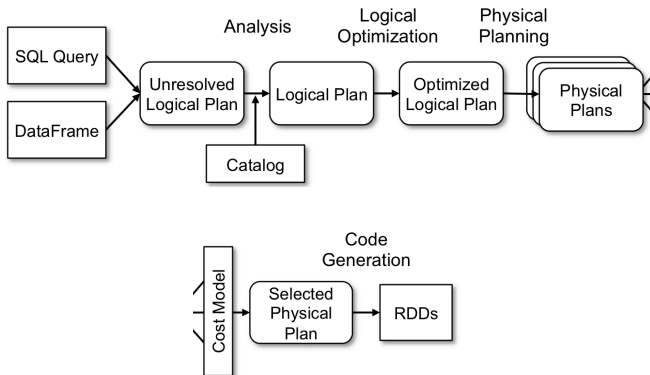


Рис.: Phases of query planning in Spark SQL. Rounded rectangles represent Catalyst trees

Example Query

```
1 query = """
2     SELECT movie_id, COUNT(*), first(title) as title
3     FROM ratings INNER JOIN movies ON ratings.movie_id == movies.movieId
4     WHERE movies.title LIKE '%(1994)%'
5     GROUP BY movie_id
6     ORDER BY COUNT(*) DESC
7 """
8
9 spark.sql(query).explain(True)
```

Example

Unresolved Logical Plan

```
1 == Parsed Logical Plan ==
2 'Sort ['COUNT(1) DESC NULLS LAST], true
3 +- 'Aggregate ['movie_id], ['movie_id, unresolvedalias('COUNT(1), None),
   first('title, false) AS title#53]
4   +- 'Filter 'movies.title LIKE %(1994)%
5     +- 'Join Inner, ('ratings.movie_id = 'movies.movieId)
6       :- 'UnresolvedRelation 'ratings'
7       +- 'UnresolvedRelation 'movies'
```

Example

Logical Plan

```
1 == Analyzed Logical Plan ==
2 movie_id: int, count(1): bigint, title: string
3 Project [movie_id#1, count(1)#56L, title#53]
4 +- Sort [count(1)#56L DESC NULLS LAST], true
5   +- Aggregate [movie_id#1], [movie_id#1, count(1) AS count(1)#56L,
6     first(title#19, false) AS title#53]
7     +- Filter title#19 LIKE %(1994)%
8       +- Join Inner, (movie_id#1 = movieId#18)
9         :- SubqueryAlias 'ratings'
10        : +- Relation[user_id#0,movie_id#1,rating#2,timestamp#3] csv
11        +- SubqueryAlias 'movies'
12          +- Relation[movieId#18,title#19,genres#20] csv
```

Example

Optimized Logical Plan

```
1 == Optimized Logical Plan ==
2 Sort [count(1)#56L DESC NULLS LAST], true
3 +- Aggregate [movie_id#1], [movie_id#1, count(1) AS count(1)#56L, first(
   title#19, false) AS title#53]
4   +- Project [movie_id#1, title#19]
5     +- Join Inner, (movie_id#1 = movieId#18)
6       :- Project [movie_id#1]
7         : +- Filter isnotnull(movie_id#1)
8         :   +- Relation[user_id#0,movie_id#1,rating#2,timestamp#3] csv
9       +- Project [movieId#18, title#19]
10        +- Filter ((isnotnull(title#19) && Contains(title#19, (1994))
11                  ) && isnotnull(movieId#18))
           +- Relation[movieId#18,title#19,genres#20] csv
```

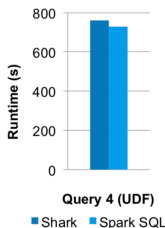
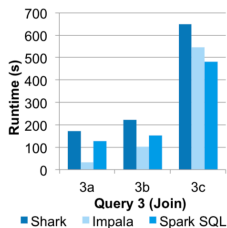
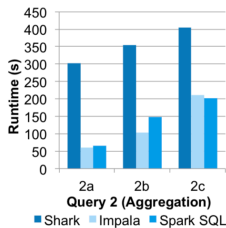
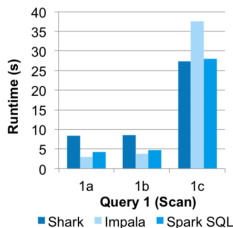
! Filter push down rule !

EVALUATION

Evaluation of the performance of Spark SQL on two dimensions:

- 1 SQL query processing performance
- 2 Spark program performance

Benchmark⁴



⁴Andrew Pavlo. "A comparison of approaches to large-scale data analysis". B: *Proceedings of the 2009 ACM SIGMOD international conference on management of data*. 2009, c. 165—178.

Distributed Aggregation

Dataset and Task

Dataset 1 billion integer pairs, (a, b) with 100000 distinct values of a

Task compute the average of b for each value of a

Distributed Aggregation

Solutions

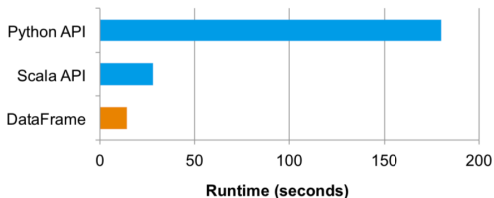
Native Spark

```
1 sum_and_count = data \  
2     .map(lambda x: (x.a, (x.b, 1))) \  
3     .reduceByKey(lambda x, y: (x[0]+y[0], x[1]+y[1])) \  
4     .collect()  
5  
6 [(x[0], x[1][0] / x[1][1]) for x in sum_and_count]
```

Spark SQL

```
1 df.groupBy("a").avg("b")
```

Distributed Aggregation Performance



In the DataFrame API, only the logical plan is constructed in Python, all physical execution is compiled into native Spark code as JVM bytecode.



Armbrust, Michael и др. "Spark sql: Relational data processing in spark". B: *Proceedings of the 2015 ACM SIGMOD international conference on management of data*. 2015, с. 1383—1394.



Karau, Holden и Rachel Warren. *High performance Spark: best practices for scaling and optimizing Apache Spark*. " O'Reilly Media, Inc.", 2017.



Kornacker, Marcel и др. "Impala: A Modern, Open-Source SQL Engine for Hadoop.". B: *Cidr*. T. 1. 2015, с. 9.



Pavlo, Andrew. "A comparison of approaches to large-scale data analysis". B: *Proceedings of the 2009 ACM SIGMOD international conference on management of data*. 2009, с. 165—178.



Xin, Reynold S и др. "Shark: SQL and rich analytics at scale". B: *Proceedings of the 2013 ACM SIGMOD International Conference on Management of data*. 2013, с. 13—24.